

# 1 Some guidance for developing new methods in BEAST 2

Disclaimer: below some ramblings on methods development for BEAST 2 (Bouckaert *et al.* 2014, Drummond and Bouckaert (2015), Bouckaert *et al.* (2019)) packages. This is a living document based on collected wisdom of BEAST developers, which keeps evolving.

This document is about testing validity of a BEAST method, not the programming aspects (like setting up dependencies, wrapping up files into a package, etc.), which can be found in the tutorial for writing a BEAST 2 package and writing a package for a tree prior tutorial.

There are several levels of validation:

- the model appears to produce reasonable results on a data set of interest.
- the model produces more reasonable results on a data set of interest than other models.
- unit tests show correctness of direct simulator implementation, likelihood implementation and/or operator implementation(s).
- sampling from prior conforms to expectations.
- a simulation study shows parameters simulated under the model can be recovered by inference from simulated data for a fixed tree and fixed other parameters for a small number of illustrative cases.
- as previous but with sampled tree and sampled parameters, so the process is repeated  $N$  times and tree and parameters sampled from a reasonable prior.
- a simulation study shows the model can recover parameters (most of the time) even when there are model violations in simulating the parameters.

Automate the experiment – you *will* do it again, in about 6 months time, when you least expect it.

Document the experiment – “It is like cleaning toilets: nobody want to do it, but it is more pleasant for visitors. You will probably be one of those visitors in 6 months time. . .”

[http://videolectures.net/cancerbioinformatics2010\\_baggerly\\_irrh/](http://videolectures.net/cancerbioinformatics2010_baggerly_irrh/)

Reproducibility with docker:

Steps to run the attached BEAST 2 analysis: 1. Install Docker ([www.docker.com](http://www.docker.com))  
2. From a terminal window, run the following from the directory containing the XML file: `docker run -v$PWD:/data tgvaughan/beast2_bacter beast ecoli.xml` That’s it! (Under Windows the \$PWD would have to be replaced with the path of the current directory.) These instructions are impervious to most things we worry about: core and package API changes, Java version changes and OS dependencies.

## 1.1 Testing new methods

New methods usually require two parts: an implementation  $I(M)$  of a model  $M$  and associated probability  $p_I(\theta|M)$  of states  $\theta$ , and MCMC operators  $R(\theta) \rightarrow \theta'$  for creating proposals  $\theta'$  for moving through state space starting in state  $\theta$  (though sometimes just an operator is validated that is much more efficient than previously existing operators). This guide contains some procedures to get some confidence that the model and operators are correctly implemented. Ideally, we have an independent implementation of a simulator  $S(M) \rightarrow \theta$  that allows (possibly inefficiently) to sample from the target distribution  $p_S(\theta|M)$ . If so, we also need to verify that the simulator is correctly implemented. In summary, we need to establish correctness of:

- the simulator implementation  $S(M) \rightarrow \theta$  (if any)
- the model implementation  $I(M)$
- operator implementations  $R$

## 2 Verify correctness of simulator implementation

To verify correctness of a simulator implementation  $S$  for model  $M$  directly, the distributions  $p_S(\theta|M)$  should match expected distribution based on theory. We can verify this by drawing a large number of samples using  $S$ , calculate summary statistics on the sample and compare these with analytical estimates for these statistics. For example, for tree priors, expected tree heights can often be determined, and for parametric distributions we often know mean and variance values. Simulating values and making sure the expected value is in the expected range is easy to verify in Tracer: the expected values should be within the mean value logged plus/minus 2 times stderr of mean (as shown in the summary statistics panel).

When no analytical estimates of statistics are available, it may be possible to find a simplified case for which an analytical solution exists, for example when the tree only has two taxa.

Examples of simulators (this list is far from exhaustive):

- the MASTER (Vaughan and Drummond 2013) BEAST 2 package is a general purpose package for simulating stochastic population dynamics models which can be expressed in terms of a chemical master equation.
- SimSnap for SNAPP (Bryant *et al.* 2012) is a custom build implementation in C++ for simulating alignments for a fixed tree and SNAPP parameters.
- The `beast.app.seqgen.SequenceSimulator` class in BEAST 2 can be used to simulate alignments for general site models using reversible substitution models. See `testSeqGen.xml` for an example.

- Models implemented in other phylogenetic software packages, such as BEAS 1, MrBayes, RevBayes, allow sampling a distribution using MCMC.
- The `beast.core.DirectSimulator` class in BEAST 2 can be used to draw samples from distributions in BEAST that extend `beast.core.distribution.Distribution` and implement the `sample(state, random)` method. You can set up an XML file and run it in BEAST. Here are a few examples: `testDirectSimulator.xml`, `testDirectSimulator2.xml`, and `testDirectSimulatorHierarchical.xml`.

### 3 Verify correctness of model implementation

For small examples for which an analytical result can be calculated a unit test can be written to confirm the implementation behaves correctly for the expected result. For example, for a small tree  $((A:1.0,B:1.0):1.0,(C:1.0,D:1.0):1.0)$  with birth rate 1 we can calculate the expected value of the Yule prior ( $\log(P) = -6$ ), and write a unit test to make sure it matches:

```
package test;

import org.junit.Test;

import beast.evolution.speciation.YuleModel;
import beast.util.TreeParser;
import junit.framework.TestCase;

public class YuleLikelihoodTest extends TestCase {

    @Test
    public void testYuleLikelihood() {
        TreeParser tree = new TreeParser("((A:1.0,B:1.0):1.0,(C:1.0,D:1.0):1.0);");

        YuleModel likelihood = new YuleModel();
        likelihood.initByName("tree", tree, "birthDiffRate", "1.0");

        assertEquals(-6.0, likelihood.calculateLogP());
    }
}
```

In theory, the inferred distributions  $p_I(\theta|M)$  should match the simulator distribution  $p_S(\theta|M)$ . However, drawing samples from  $p_I(\theta|M)$  typically requires running an MCMC chain, which requires MCMC proposals  $R$  to randomly walk through state space. If we do this, we need to rely on  $R$  being correctly implemented. So, if we find that  $p_I(\theta|M)$  and  $p_S(\theta|M)$  do not match, it is not possible to tell whether problem is with an operator  $R$  or with the model

implementation  $I(M)$ .

The Hastings ratio for this operator is  $P(\theta)/P(\theta')$ . Consequently, every proposal is accepted, whether  $p_I(\theta|M)$  is correctly implemented or not.

In BEAST, if the `sample` method is implemented in a class derived from `Distribution`, you can use `beast.experiment.DirectSimulatorOperator` in the `Experimenter` package to set up an MCMC analysis in XML. Here is an example that draws a birth rate from an exponential distribution with mean 1, and a Yule distribution to generate a tree. Note that the tree height statistic is logged, as well as an expression for a clock rate (being  $0.5/\text{tree-height}$ ) for evaluation purposes. The MCMC sample can be compared with the direct sample using the example file `testDirectSimulator.xml`.

```
<beast version="2.0" namespace="beast.core
:beast.evolution.alignment
:beast.evolution.tree
:beast.math.distributions
:beast.evolution.speciation
:beast.core.util
:beast.core.parameter">

  <run spec="MCMC" chainLength="1000000">
    <state id="state">
      <stateNode idref="tree"/>
      <stateNode idref="birthDiffRateParam"/>
    </state>

    <distribution spec="CompoundDistribution" id="fullModel">
      <distribution spec="YuleModel" id="yuleModel">
        <tree spec="Tree" id="tree">
          <taxonset spec="TaxonSet">
            <taxon spec="Taxon" id="t1"/>
            <taxon spec="Taxon" id="t2"/>
            <taxon spec="Taxon" id="t3"/>
            <taxon spec="Taxon" id="t4"/>
            <taxon spec="Taxon" id="t5"/>
          </taxonset>
        </tree>
        <birthDiffRate spec="RealParameter" id="birthDiffRateParam" value="1.0"/>
      </distribution>

      <distribution spec="beast.math.distributions.Prior" id="birthDiffRatePrior">
        <distr spec="Exponential" id="xExpParamDist" mean="1"/>
        <x idref="birthDiffRateParam"/>
      </distribution>
    </run>
  </beast>
```

```

        </distribution>

</distribution>

<operator spec="beast.experiment.DirectSimulatorOperator" weight="1" state="@state">
    <simulator id="DirectSimulator" spec="beast.core.DirectSimulator" nSamples="1">
        <distribution idref="fullModel"/>
    </simulator>
</operator>

<logger id="tracelog" logEvery="1000" fileName="$(filebase).log">
    <log idref="birthDiffRateParam"/>
    <log id="clockRate" spec="beast.util.Script" expression="0.5/TreeHeight">
        <x id="TreeHeight" spec="beast.evolution.tree.TreeHeightLogger" tree="@tree">
        </x>
    </log>
    <log idref="TreeHeight"/>
</logger>

<logger id="treelog" logEvery="1000" fileName="$(filebase).trees">
    <log idref="tree"/>
</logger>

<logger id="screenlog" logEvery="1000">
    <log idref="birthDiffRateParam"/>
</logger>
</run>
</beast>

```

Make sure when sampling from the prior through MCMC that the chain length is sufficiently large and log frequency large enough to ensure that each sample is independent of the previous sample. The ESSs shown in Tracer should be close to N when there are N samples in the trace log, for most of the items in the log. There may be a few items with an ESS that is a bit lower, and inspection of the trace plot should tell you whether lower log frequencies should be used.

Comparing two distributions can be done by

- eye balling the marginal likelihoods in Tracer and making sure they are close enough.
- testing whether parameters are covered 95% of the time in the 95% HPD interval of parameter distributions.
- using a statistical test, e.g. the Kolmogorov-Smirnov test, to verify the distributions  $p_I(\theta|M)$  and  $p_S(\theta|M)$  are the same.

**TraceKSStats** calculate Kolmogorov-Smirnof statistic for comparing trace logs. **TraceKSStats** has the following inputs:

- trace1 (LogFile): first trace file to compare (required)

- trace2 (LogFile): second trace file to compare (required)
- burnin (Integer): percentage of trace logs to used as burn-in (and will be ignored) (optional, default: 10)

Sample output:

Trace entry	p-value
posterior	1.0
likelihood	0.21107622404022763
prior	0.036794035181748064
treeLikelihood	0.04781117967724258
TreeHeight	0.036794035181748064
YuleModel	0.005399806065857771
birthRate	0.2815361702146215
kappa	0.62072545444263
freqParameter.1	0.0
freqParameter.2	8.930072172019798E-5
freqParameter.3	1.0883734952171764E-6
freqParameter.4	0.0

Though some values have very low p-values, meaning they differ significantly, it is recommended to verify this using Tracer to make sure that the test is not unduly influenced by outliers.

### 3.1 Score based model validation

For a density  $p$  on data  $x$  (this could be a scalar, vector, or a tree) parameterised by the vector  $\theta$ , the expected value of the score function  $U(\theta, x) = \frac{\partial}{\partial \theta} \log p(x; \theta)$  is 0:

$$E(U(\theta, x)) = \int U(\theta, x)p(x; \theta)dx = 0 \quad (1)$$

(this is the expected value over the data  $x$  at the true parameters  $\theta$ .) Also, the covariance of the score function is equal to the negative Hessian of the log-likelihood:

$$E(U(\theta, x)U(\theta, x)^T + \partial^2 / \partial \theta^2 \log p(x; \theta)) = 0 \quad (2)$$

These properties can be used in conjunction with a direct simulator as a necessary but not sufficient check that the likelihood is implemented correctly. The score function (eq. 1) and Hessian (eq. 2) statistics can be calculated on samples from the simulator and a hypothesis test used to check that their mean is 0. In the multivariate case a potentially useful test is the likelihood ratio test for a multivariate normal with zero mean (implemented here). If there are

non-identifiable parameters there may be issues with performing this test as colinearity will lead to a singular sample covariance matrix.

In practice it is appropriate to compute the score function by calculating derivatives using finite differences (implemented here).

### 3.1.1 Writing a test

The BEAST validation package is designed around three core object types: it performs a *test* on *statistics* drawn from *samplers*. BEAST validation tests are implemented within the BEAST 2 XML parser framework. The main Runnable class is `beast.validation.StochasticValidationTest`. `StochasticValidationTest` has inputs for each of the core object types: `samplers`, `statistics` and a `test`. Note that some tests may be designed for one, two or many sampler/statistic pairs.

`StochasticValidationTest` has some additional parameters:

- `alpha`: The significant level to use in the test
- `nSamples`: The number of samples to draw from each sampler
- `printEvery`: How often to report sampling progress
- `sampleLoggers`: Loggers to run for every sample (usually the statistics)
- `resultLoggers`: Loggers to run after testing

There are currently two useful combinations of samplers, statistics and test implemented by BEAST validation.

### 3.1.2 Score function validation

This test validates a combination of likelihood and direct simulator using a known property of probability density functions: that the expectation of the gradient of the log-likelihood at the true parameter values is zero (see `stats-tricks.md`). The core components of this test are (with a single sampler-statistic pair):

- a simulator: This could be a custom simulator, or make use of one of the generic simulation tools available, such as `beast.simulation.TreeSamplerFromMaster`
- `beast.validation.statistics.NumericalScoreFunctionStatistics`: a statistic that uses a finite differences to calculate the gradient of a likelihood with respect to some parameters
  - Note that the `RealParameter` objects included in the `parameter` input should be the same provided to the `Distribution` in the `likelihood` input
- `beast.validation.tests.MultivariateNormalZeroMeanTest`: a likelihood ratio test that uses a multivariate normal fit to the gradients to test for zero mean

An important note is that there are some regularity conditions on the parameters you can use in this test. Roughly, they must not affect the support of the data, which excludes the origin time parameter in tree priors. See `stats-tricks.md` for further details.

An example XML for this test on the birth-death-sampling tree prior can be found in the BEAST validation examples.

### 3.1.3 Running a test

Once you have created an XML and installed the BEAST validation package (pending inclusion in the package repository) using the standard BEAST 2 launcher. Once your test has run, it will provide the result on the console:

```
Performing test...
Test PASSED
p value: 0.284784
```

## 4 Verify correctness of operator implementations

Once simulator  $S(M)$  and model implementation  $I(M)$  are verified to be correct, next step is implementing efficient operators, running MCMCs to verify that parameters drawn from the prior are covered 95% of the time in the 95% HPD interval of parameter distributions.

The direct simulator operator (see `DirectSimulatorOperator` above) can be used as starting operator, and new operators added one by one to verify correctness.

The BEAST 2 Experimenter package can assist (see section “Using the Experimenter package” below).

### 4.1 MCMC sampling to verify correctness of operator implementations

A useful test for an MCMC sampler for a model (likelihood + operators) is if it can produce the same distribution as direct simulation. For phylogenetic models, this would usually be the tree prior. The core components of this test:

- Two samplers
  - a simulator (see the previous test for details)
  - `beast.core.SamplerFromMCMC`
    - \* extends the normal BEAST MCMC class



- \* needs a model/likelihood and operators
- \* To test a single, potentially non-ergodic operator, a separate simulator could be used as a global operator with the `beast.simulation.OperatorFromSampler` class
- A multivariate statistic
  - `beast.validation.statistics.UltrametricTreeStatistics` provides some basic statistics on BEAST time trees
  - For an example of statistics on more complex tree-like objects, see `bacter.util.ConversionGraphStatsLogger`
- `beast.validation.tests.BootstrapMultivariateDistributionTest`: a test that bootstraps a multivariate generalisation of the Kolmogorov-Smirnov (KS) statistic to compare distributions (see [stats-tricks.md] for further details)

An example XML for this test on the birth-death-sampling tree prior can be found in the BEAST validation examples.

#### 4.1.1 The dreaded “likelihood incorrectly calculated” error message

If you see this error message during an MCMC run, there probably is an error in the way that a `CalculationNode` in your model is caching its state. More information is on the [beast site blog post](#).

## 4.2 Setting up a direct simulation in BEAST

Using the direct simulator can be done as follows

- Set up `DirectSimulator` at top level
- Add model to simulate from
- Add loggers to register output

#### 4.2.1 Set up `DirectSimulator` at top level

Use the `nSamples` attribute to specify how many samples to draw.

```
<beast version="2.0" namespace="beast.core:beast.evolution.alignment:beast.evolution.tree:beast.evolution.simulation">
  <run spec="DirectSimulator" nSamples="100">
    <!-- model goes here -->
    <!-- loggers go here -->
  </run>
</beast>
```

### 4.2.2 Add model to simulate from

Here, the model consists of a `birthDiffRate` parameter, drawn from an exponential distribution. This rate is used in a Yule model to draw trees over a set of 5 taxa, called `t1`, `t2`, ..., `t4`. This is placed inside the `run` element above.

```
<distribution spec="CompoundDistribution" id="fullModel">
  <distribution spec="YuleModel" id="yuleModel">
    <tree spec="Tree" id="tree">
      <taxonset spec="TaxonSet">
        <taxon spec="Taxon" id="t1"/>
        <taxon spec="Taxon" id="t2"/>
        <taxon spec="Taxon" id="t3"/>
        <taxon spec="Taxon" id="t4"/>
        <taxon spec="Taxon" id="t5"/>
      </taxonset>
    </tree>
    <birthDiffRate spec="RealParameter" id="birthDiffRateParam" value="1.0"/>
  </distribution>

  <distribution spec="beast.math.distributions.Prior" id="birthDiffRatePrior">
    <distr spec="Exponential" id="xExpParamDist" mean="1"/>
    <x idref="birthDiffRateParam"/>
  </distribution>
</distribution>
```

### 4.2.3 Add loggers to register output

Since we sample a parameter and a tree, we need a trace log and a tree log. Apart from state nodes, other statistics can be sampled as well. For example, here, we log the height of the tree using a `TreeHeightLogger`, and log a clock rate suitable for the tree using the expression `0.5/TreeHeight` using the `Script` class from the BEASTLabs package. The loggers should be placed inside the `run` element as well.

```
<logger logEvery="1" fileName="$(filebase).log">
  <log idref="birthDiffRateParam"/>
  <log id="TreeHeight" spec="beast.evolution.tree.TreeHeightLogger" tree="@tree"/>
  <log id="clockRate" spec="beast.util.Script" expression="0.5/TreeHeight">
    <x idref="TreeHeight"/>
  </log>
</logger>

<logger logEvery="1" fileName="$(filebase).trees">
  <log idref="tree"/>
</logger>
```

The complete XML file can be found as `testDirectSimulatorByMCMC.xml` in the `Experimenter` package.

### 4.3 Converting direct simulator XML to MCMC

Conversion requires the following steps:

- replace top level run element by MCMC
- add state element and references to state nodes
- add `DirectSimulatorOperator`
- add screen logger (optional)

### 4.4 Dealing with properties of branches in trees

When you have meta data associated with branches of trees, like population sizes or clock rates, it is usually easy to verify that the meta data that is inferred is compatible for the leaf nodes. Leaf nodes in a tree with  $n$  taxa are numbered  $0, \dots, n - 1$ , and at the start of a BEAST run these numbers are assigned and do not change during the MCMC execution. So, if you have a parameter associated with branches, the first  $n$  values will be associated with the leaf branches, and can be logged in a `tracelog` and compared with values used to generate the data, just like any other parameter.

However, the branches that do not end in a leaf are numbered  $n, \dots, 2n - 2$  with potentially root branch  $2n - 1$ . Furthermore, branches are not associated with clades, and their numbers can change throughout the MCMC run. This means that comparing value  $n$  of a parameter with the logged value at dimension  $n$  may not match the correct values.

One way to verify the correctness of estimates of metadata on branches is do a well calibrated simulation study that

- logs the tree + metadata used to generate the alignment.
- create a summary tree with topology equal to the generating tree.
- run `MCCTreeComparator` to summarise results.

#### 4.4.1 1. logging generating tree + metadata

In the BEAST XML template, add a tree logger that logs the tree used to generate the alignment, e.g like so:

```
<logger id="sim.treelog.t:alignment" spec="Logger" fileName="sim.%(filebase).trees" log
  <log id="sim.TreeWithMetaDataLogger.t:alignment" spec="beast.phoneme.BranchMetadData
    <metadata idref="sim.permutation.s:alignment"/>
  </log>
</logger>
```

We only need this to be logged once, to the `logEvery` value can be set very high.

#### 4.4.2 2. create a summary tree with topology equal to the generating tree

Say, we stored the generating trees in files `sim.analysis-0.trees`, `sim.analysis-1.trees`, etc. numbered 0 through to 99. Treeannotator (part of BEAST) can be used to summarise trees and their metadata on the tree topology used to generate the data using the `-target` option. On unix-like systems, the following command will generate summary trees for each of the 100 runs:

```
for i in {0..99}; do
    treeannotator -b 10 -target sim.analysis-$i.trees analysis-$i.trees analysis-$i.tree
done
```

#### 4.4.3 3. run MCCTreeComparator to summarise results.

MCCTreeComparator is a tool that is part of the beast-validation package, which produces coverage summaries for branch metadata.

MCCTreeComparator has the following options:

- `tree` (TreeFile): source tree file with meta data (required)
- `mcc` (TreeFile): MCC source tree file (required)
- `from` (Integer): start value to loop over (optional, default: 0)
- `to` (Integer): end value (inclusive) to loop over. If less than 0, no loop is performed. If more than 0, the part `$(n)` in the file path will be replaced by an integer, starting at 'from' and ending in 'to' (optional, default: 99)
- `out` (OutFile): output file, or stdout if not specified (optional, default: `[[none]]`)

The output is a small report containing coverage information for total, leaf branches and internal branches separately. For real valued metadata it is the percentage of clades where the generating value is in the 95% HPD that is estimated (like clade height). For categorical values (below p1, p2, p3 and p4) it is the percentage of clades where the generating value is in the 95% credible set. The `height` and `posterior` entries show an 'N/A' for not applicable for leaf branch matches.

metadata	% matches	% leaf branch matches	% internal branch matches
height	95	N/A	95
p1	99.3	99.5	99
p2	99.6	99.83	99.25
p3	99.9	100	99.75
p4	99.1	99.83	98

metadata	% matches	% leaf branch matches	% internal branch matches
posterior	99.6	N/A	99.6

## 5 Practical considerations

Validation only covers cases in as far as the prior covers it – most studies will not cover all possible cases, since the state space is just too large. Usually, informative priors are required for validation to work, since broader priors (e.g. some of the default tree priors in BEAST) lead to identifiability issues, for example, due to saturation of mutations along branches of a tree.

### 5.1 Setting priors

#### 5.1.1 Trees & clock model parameters

The mutation rate  $\mu$  must have been such that the tree height  $h_T$  cannot exceed  $1/\mu$  (in other words,  $\mu h_T \leq 1$ ), otherwise there would be saturation, and sequences could not possibly have sufficient information to align. At the other end of the spectrum, where  $\mu h_T$  close to zero, very long sequences are required to ensure there are enough mutations in order to be able to reconstruct the tree distribution.

- for reasonable computation times, trees should be about 0.5 substitutions high, OR
- sequences should be very long to reliably reconstruct smaller trees.

One way to enforce this is by

- a narrow prior on birth rates (for birth/death type tree priors), or
  - putting an MRCA prior on the height of the tree, for coalescent models.
- Note that the latter hampers direct simulator implementations.

For clock models with mean clock rate  $\neq 1$ , simulate trees with clock rates times tree height approximately 0.5.

Published mutation rates can range from  $O(1e - 2)$  substitutions per site per year for viruses such as HIV (Cuevas *et al.* 2015), to  $O(1e - 11)$  for conserved regions of nuclear DNA (e.g. PyrE2 locus in *Haloferax volcanii* (Lynch 2010)).

For releases, tree priors for clock and trees should be made less informative in order to cater for a wider range of tree heights and clock rates.

### 5.1.2 Gamma rate heterogeneity

To prevent saturation, adding categories with slow rates will go some way to allow covering a larger range of clock rates. Using gamma rate heterogeneity with shape values in the range 0.1 to 1 allows this, so adopt a gamma shape prior accordingly.

### 5.1.3 Proportion invariable sites

Since each site evolves with non-zero rate, use of proportion invariable sites is modelling the process badly, and therefore not recommended.

### 5.1.4 Frequencies

Priors ideally should be set in realistic ranges, e.g. frequency priors not uniform(0,1) but Dirichlet(4,4,4,4) is better.

### 5.1.5 Substitution model parameters

Default priors seem OK for most substitution models.

## 5.2 Sequence simulator

`SequenceSimulator` (alternatively `SimulatedAlignment`) can help generate individual alignments.

- `SequenceSimulator` can be used to generate a static alignment that can be merged into existing XML (e.g. with the help of `beast.app.seqgen.MergeDataWith`).
- `SimulatedAlignment` can be used as replacement of an alignment in the XML and will generate a new alignment every time BEAST is started on the XML.

To generate N XML files, use `CoverageTestXMLGenerator` in `Experimenter` package

The sequence length should be long enough that trees can be reasonably reliably recovered – if the difference between longest and shorted tree is 2 orders of magnitude, nucleotide sequences of 10 thousand sites. When  $\mu$ \*tree-height approximate 0.5, sequences of length 1000 are sufficient.

### 5.3 Log file names

Make sure log files names do not overlap. Use `logFileName="out$(N).log` and start BEAST with `for i in {0..99} do /path/to/beast/bin/beast -D N=$i beast$i.xml; done`

## 6 Trouble shooting

### 6.1 Coverage gone wrong

One reason coverage can be lower is if the ESSs are too small, which can be easily checked by looking at the minimum ESS for the log entry. If these values are much below 200 the chain length should be increased to be sure any low coverage is not due to insufficient convergence of the MCMC chain.

### 6.2 Model misspecification by using empirical estimates (e.g. frequencies)

When using empirical frequencies, these frequencies can bias estimates of other parameters (like gamma shape for gamma rate heterogeneity), causing low coverage for these parameters. Especially for short sequences, empirical estimates can be far away from the frequencies used to generate the data. For that reason, empirical frequencies should be avoided, and estimated frequencies used instead.

Another argument against empirical frequencies is that it is double dipping: using the data to both estimate frequencies and ...

### 6.3 Low coverage

The occasional 91 is acceptable (the 95% HPD = 90 to 98 probability the implementation is correct) but coverage below 90 almost surely indicate an issue with the model or operator implementation. Also, coverage of 99 or 100 should be looked at with suspicion – it may indicate overly wide uncertainty intervals.

If correct, distributed binomial with  $p=0.95$ ,  $N=100$ :

k	$p(x=k)$	$p(x \leq k)$	$p(x > k)$
90	0.0167	0.0282	0.9718
91	0.0349	0.0631	0.9369
92	0.0649	0.1280	0.8720
93	0.1060	0.2340	0.7660
94	0.1500	0.3840	0.6160

k	$p(x=k)$	$p(x \leq k)$	$p(x \geq k)$
95	0.1800	0.5640	0.4360
96	0.1781	0.7422	0.2578
97	0.1396	0.8817	0.1183
98	0.0812	0.9629	0.0371
99	0.0312	0.9941	0.0059
100	0.0059	1.0000	0.0000

source <https://www.di-mgt.com.au/binomial-calculator.html> for different values of  $p$  and  $N$ .

## 6.4 Common causes of low coverage

- ESS too low
- `loganalyser` was called with `out*.log` instead of `out?.log` `out???.log` `out????.log`. As a result, the `loganalyser` output and trace log of true values are in different order. Rerun `loganalyser` to make sure the output is in the right order.
- improper priors used: all priors should be proper, that is integrate to 1. Examples of improper priors are the  $1/X$  and uniform prior with infinite upper and/or lower bounds.
- priors are outside the range usually used in applying the model, especially the next case:
- trees cannot be reconstructed reliably (height should not be too small or large).
- Hastings ratio in operators incorrectly implemented
- bug in model likelihood

## 7 Releasing

### 7.1 Setting priors in BEAUti template

- For released software, default priors should be made as uninformative as possible and/or provide some guidance on how to set them – users will use defaults.
- Make sure to consider a number of scenarios, e.g. for clock models, scenarios from setting up rates page on [beast2.org](http://beast2.org).



## 7.2 Communicating results

When publishing well calibrated studies, all XML files, log files and scripts for manipulating them should be made available, so the study can be replicated exactly as is. A practical way to do this is through a github repository.

Version numbers of BEAST and packages used should be noted.

## 8 Using the Experimenter package

Experimenter is a BEAST 2 package that assists in simulation studies to verify correctness of the implementation. The goal of this particular simulation studies is to make sure that the model or operator implementation is correct by running N analysis on simulated data (using SequenceSimulator) on a tree and site model parameters sampled from a prior.

To run a simulation study:

- set up XML for desired model and sample from prior
- generate (MCMC) analysis for each of the samples (say 100)
- run the analyses
- use loganalyser to summarise trace files
- run `CoverageCalculator` to summarise coverage of parameters

Make sure to have the Experimenter package installed (details at the end).

### 8.1 Step 1. Set up XML for desired model and sample from prior

First, you set up a BEAST analysis in an XML file in the configuration that you want to test. If you are using MCMC, set the `sampleFromPrior="true"` flag on the element with MCMC in it, and sample from the prior.

Alternatively, use a `DirectSimulator` as in this example: `testDirectSimulator.xml`.

Make sure that the number of samples in the trace log and tree log is the same and that they are sampled at a frequency such that there will be N useful samples (say N=100).

### 8.2 Step 2. Generate (MCMC) analysis for each of the samples (with single alignment)

`CoverageTestXMLGenerator2` Generate XML for performing coverage test (using `CoverageCalculator`). A template XML file should contain

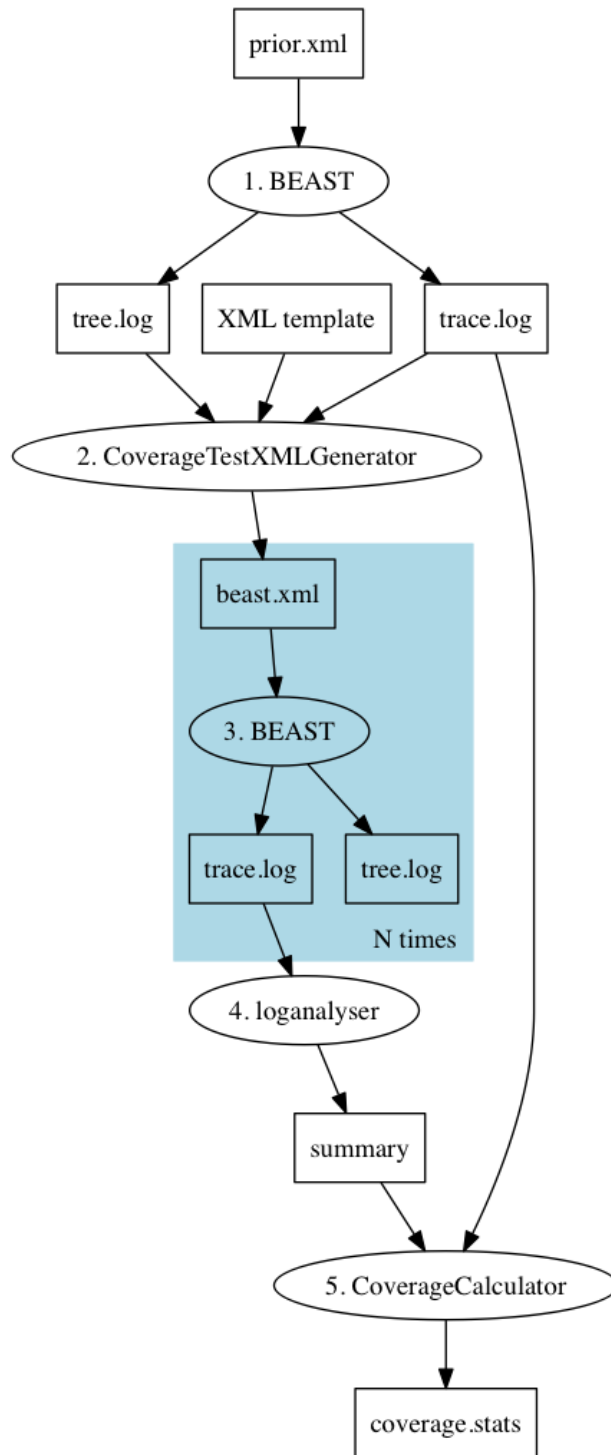


Figure 1: Summary of files involved in testing an operator. Rectangles represent files, ovals represent programs.

**SimulatedAlignments** with a tree, branch rate model and site model specified. These can be parameterised by using string of the form `$(n)` where `n` should match exactly the columns in the trace log file, or it can be `$(tree)` which is replaced by the species tree as specified in the `-treeFile` option. For individual gene trees use the `-geneTreeFile` option (detailed below). Example input files can be found [here](#).

It has the following options:

- `workingDir` (File): working directory where input files live and output directory is created (optional)
- `outDir` (String): output directory where generated XML goes (as sub dir of working dir) (optional, default: `mcmc`)
- `logFile` (LogFile): trace log file containing model parameter values to use for generating sequence data (optional)
- `treeFile` (TreeFile): tree log file containing trees to generate sequence data on (optional)
- `geneTreeFile` (File): (optional) configuration file with gene tree identifiers and log file names, one per line separated by a tab, for example

```
gene1    /xyz/abc/gene1.trees
gene2    /xyz/abc/gene2.trees
```

In the XML, any instance of `$(gene1)` will be replaced by a tree from the file `/xyz/abc/gene1.trees`, and likewise for `$(gene2)`.

- `xmlFile` (XMLFile): XML template file containing analysis to be merged with generated sequence data (optional)
- `skip` (Integer): numer of log file lines to skip (optional, default: 1)
- `help` show arguments

NB: make sure to set `sampleFromPrior="false"` in the XML.

NB: to ensure unique file name, add a parameter to `logFileName`, e.g. `logFileName="out$(N).log"`

With this setting, when you run BEAST with `-D N=1` the log file will be `out1.log`.

NB each **SimulatedAlignment** should have its own tree, `sitemodel` and branch rate model, and sh

### 8.3 Alternative Step 2. Generate (MCMC) analysis for each of the samples (support for multiple alignments)

Instead of generating complete BEAST XML files, including alignments, XML files can be generated where the state is initialised by start values from the

trace log and the tree is initialised with trees from the tree log in Newick format through `TreeParser`. The alignment can then be generated on the fly with the help of `beast.app.seqgen.SimulatedAlignment`. Note that every time BEAST is started on the XML a new alignment will be generated, and unless the same seed is used, these alignments will be unique (almost surely).

## 8.4 Step 3. Run the analyses

Use your favourite method to run the  $N$  analyses, for example with a shell script

```
for i in {0..99} do /path/to/beast/bin/beast -D N=$i beast$i.xml; done
```

where `/path/to/beast` the path to where BEAST is installed.

## 8.5 Step 4. Use loganalyser to summarise trace files

Use the loganalyser utility that comes with BEAST in the bin directory. It is important to use the `-oneline` argument so that each log line gets summarised on a single line, which is what `CoverageCalculator` expects. Also, it is important that the log lines are in the same order as the log lines in the sample from the prior, so put the results for single digits before those of double digits, e.g. like so:

```
/path/to/beast/bin/loganalyser -oneline out?.log out??.log > results
```

where `out` the base name of your output log file.

NB if loganalyser is called with ``out*.log``, results may end up in order ``1,10,11,...,19,2,2``

## 8.6 Step 5. Run CoverageCalculator to summarise coverage of parameters

You can run `CoverageCalculator` using the BEAST applauncher utility (or via the `File/Launch Apps` meny in BEAUti).

`CoverageCalculator` calculates how many times entries in log file are covered in an estimated 95% HPD interval and has the following arguments:

- `log <filename>` log file containing actual values
- `skip <integer>` numer of log file lines to skip (default: 1)
- `logAnalyser <filename>` file produced by loganalyser tool using the `-oneline` option, containing estimated values
- `out <directory>` output directory for tsv files with truth and estimated mean and 95% HPDs, and directory is also used to generate svg bargraphs and html report. Not produced if not specified.

- `typeFile` (File): if specified, the type file is a tab delimited file with first column containing entry names as they appear in the trace log file, and second column variable type, d for double, b for binary, c for categorical, for example:

<code>variable</code>	<code>type</code>
<code>birthRate</code>	d
<code>kappa</code>	d
<code>hasGamma</code>	b
<code>modelIndicator</code>	c

Items that are not specified are considered to be of type double.

It produces a report like so:

	coverage	Mean ESS	Min ESS
posterior	0	2188.41	1363.02
likelihood	0	4333.99	3042.15
prior	33	1613.20	891.92
treeLikelihood.dna	0	4333.99	3042.15
TreeHeight	95	3076.44	2233.29
popSize	94	577.20	331.78
CoalescentConstant	91	1620.76	787.30
logP(mrca(root))	97	4320.70	3328.88
mrca.age(root)	95	3076.44	2233.29
clockRate	0	3046.64	2174.60
freqParameter.1	98	4332.76	3388.90
freqParameter.2	97	4337.93	3334.29
freqParameter.3	96	4378.30	3462.73
freqParameter.4	92	4348.83	3316.36

Coverage should be around 95%. One reason coverage can be lower is if the ESSs are too small, which can be easily checked by looking at the **Mean ESS** and **Min ESS** columns. If these values are much below 200 the chain length should be increased to be sure any low coverage is not due to insufficient convergence of the MCMC chain. The occasional 90 or 91 is acceptable but coverage below 90 almost surely indicate an issue with the model or operator implementation.

The values for posterior, prior and treelikelihood can be ignored: it compares results from sampling from the prior with that of sampling from the posterior so they can be expected to be different.

If an output file is specified, **CoverageCalculator** also generates an HTML file with bar graphs (in `svg` ) showing how well each item in the log file covers the true value, as well as tab separated (`tsv`) files containing the data, so you can import them in for example R to produce customised graphs. Below some examples with good coverage, and strong, medium and weak ability to learn the parameter, followed by over estimated and under estimated parameters.

Graphs show true value on x-axis and estimates on y-axis. Black line shows where




Figure 2: Strong coverage: true parameter can be inferred accurately.




Figure 3: Medium coverage: true parameter can be inferred, but with high uncertainty.




Figure 4: Weak coverage: true parameter cannot be inferred, even though 95% HPD covers the true value sufficiently often.




Figure 5: True parameter is over estimated




Figure 6: True parameter is under estimated

x equals y axis, and where ideally most of the probability mass is concentrated. Black dots are means of estimates. Bars indicate 95% HPDs where blue bars cover the true value and red ones do not. Ideally 95 out of 100 bars should be blue.

## 8.7 Simulation Based Calibration

Simulation Based Calibration (SBC) (Talts *et al.* 2018) is a way to validate how well true values used to generate data rank inside the inferred distributions. Ranks are binned, and the resulting bins should be uniformly distributed if all is well. Deviation from uniform distributions indicate

- if shaped like a U the posterior is too narrow.
- if shape like inverted U, the posterior is too wide.
- if shaped sloping upwards, the posterior is biased towards lower estimates.
- if shaped sloping downwards, the posterior is biased towards higher estimates.

To run a simulated based calibration study (steps 1-3 as for a coverage study):

- set up XML for desired model and sample from prior
- generate (MCMC) analysis for each of the samples (say 100)
- run the analyses
- use **LogAnalyser** to find minimum ESS






Figure 7: Simulation based calibration output for kappa parameter with 20 intervals. Light grey band indicates 99% coverage interval and darker grey the 95% coverage interval.

- run **LogCombiner** to sub sample log files and accumulate logs
- run **SBCAnalyser** to summarise coverage of parameters

A correct implementation is uniformly distributed, like so:

For steps 1-3, see coverage study.

## 8.8 Step 4. Find minimum ESS

Use **LogAnalyser** to find minimum ESS – or run coverage study and minimum ESS will be printed as part of the analysis.

## 8.9 Step 5. Combine logs

First, we need to determine how much to resample log files. Since samples must be independent for the method to work, we can resample with frequency equal to the chain length divided by minimum ESS.

Run **LogCombiner** to sub sample log files and accumulate logs. To run from command line, use

```
/path/to/beast/bin/logcombiner -resample <resample> -log <name>-?.log <name>-???.log <name>-?
```

where **<resample>** is the resample frequency (= chain length/minimum ESS), and **<name>-** the name of the log file. Note that if you numbered the log files 0, ..., 9, 10, ..., 99, 100, ..., 999 using **<name>-\*?.log** will put entries in an alphabetic order, which is probably *not* what you want.

## 8.10 Step 6. Run SBCAnalyser to summarise coverage of parameters

**SBCAnalyser** can be run with the BEAST app launcher, and outputs a report and (if an output directory is specified). It has the following arguments:

- **SBCAnalyser** has the following inputs:
- **log <filename>**: log file containing actual values (required)
- **skip <integer>**: number of log file lines to skip (optional, default: 1)
- **logAnalyser <filename>**: file produced by loganalyser tool using the **-online** option, containing estimated values (required)
- **bins <integer>**: number of bins to represent prior distribution. If not specified (or not positive) use number of samples from posterior + 1 ( $L+1$  in the paper) (optional, default: -1)
- **outputDir <directory>**: output directory for SVG bar charts (optional, default: `[[none]]`)
- **useRankedBins <boolean>**: if true use ranking wrt prior to find bins. If false, use empirical bins based on prior. (optional, default: true)

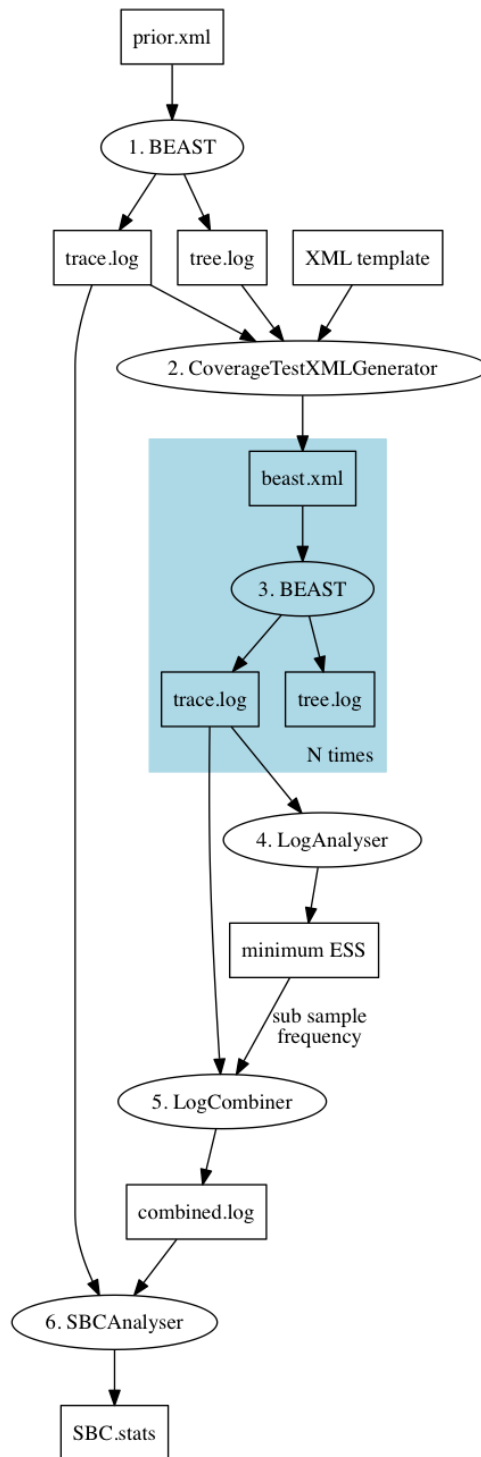


Figure 8: Summary of files involved in running a simulation based calibration study. Rectangles represent files, ovals represent programs.

Note that it compares entries from the prior to posterior, so items like likelihood, posterior, treeLikelihood and clockRate seem very wrong, but that can be ignored, since these were not part of the prior or (for clockRate) we know beforehand the prior differs substantially from the posterior.

99%lo << mean << 99%up = -1 << 5 << 10

	missed	bin0		bin1		bin2		bin3	
posterior	1	0	0	0	0	0	0	0	0
likelihood	1	0	0	0	0	0	0	0	0
prior	0	8	5	6	6	4	3	7	5
treeLikelihood.dna	1	0	0	0	0	0	0	0	0
TreeHeight	1	3	5	3	3	2	4	5	3
kappa	0	8	3	5	5	5	4	4	6
gammaShape	0	6	6	2	4	3	2	2	6
popSize	1	1	3	1	3	2	6	1	3
CoalescentConstant	0	7	5	0	7	9	5	5	4
parameter.hyperInverseGamma-beta-PopSizePrior	0	6	3	2	6	3	3	3	7
HyperPrior.hyperInverseGamma-beta-PopSizePrior	1	11	2	5	4	5	3	5	10
monophyletic(root)	1	0	0	0	0	0	0	0	0
logP(mrca(root))	0	5	5	1	4	2	4	6	5
mrca.age(root)	1	3	5	3	3	2	4	5	3
clockRate	1	0	0	0	0	0	0	0	0
freqParameter.1	1	4	4	7	3	6	7	4	6
freqParameter.2	0	6	9	3	5	4	7	4	4
freqParameter.3	0	6	4	6	8	8	5	4	5
freqParameter.4	0	9	1	2	5	4	3	4	6
Done!									

## 8.11 Installing Experimenter package

Currently, you need to build from source (which depends on BEAST 2, BEASTlabs and MASTER code) and install by hand (see “install by hand” section in managing packages).

Quick guide

- clone BEAST 2, BEASTlabs, MASTER, and Experimenter all in same directory.
- build BEAST 2 (using `ant Linux` in the `beast2` folder), then BEASTLabs (using `ant addon` in the `BEASTLabs` folder), and MASTER (using `ant` in the `MASTER` folder) then Experimenter (again, using `ant addon` in the `Experimenter` folder) packages.
- install BEASTlabs (using the package manager, or via BEAUti’s `File/Manage packages` menu).
- install Experimenter package by creating `Experimenter` folder in your BEAST package folder, and unzip the file `Experimenter/build/dist/Experimenter.addon.v0.0.1.zip`

(assuming version 0.0.1).

## 8.12 References

- Bouckaert R, Vaughan TG, Barido-Sottani J *et al.* BEAST 2.5: An advanced software platform for bayesian evolutionary analysis. *PLoS computational biology* 2019;**15**:e1006650.
- Bouckaert RR, Heled J, Kühnert D *et al.* BEAST 2: A software platform for Bayesian evolutionary analysis. *PLoS Comput Biol* 2014;**10**:e1003537.
- Bryant D, Bouckaert R, Felsenstein J *et al.* Inferring species trees directly from biallelic genetic markers: Bypassing gene trees in a full coalescent analysis. *Molecular biology and evolution* 2012;**29**:1917–32.
- Cuevas JM, Geller R, Garijo R *et al.* Extremely high mutation rate of hiv-1 in vivo. *PLoS biology* 2015;**13**:e1002251.
- Drummond AJ, Bouckaert RR. *Bayesian Evolutionary Analysis with BEAST*. Cambridge: Cambridge University Press, 2015.
- Lynch M. Evolution of the mutation rate. *TRENDS in Genetics* 2010;**26**:345–52.
- Talts S, Betancourt M, Simpson D *et al.* Validating bayesian inference algorithms with simulation-based calibration. *arXiv preprint arXiv:1804.06788* 2018.
- Vaughan TG, Drummond AJ. A stochastic simulator of birth–death master equations with application to phylodynamics. *Molecular biology and evolution* 2013;**30**:1480–93.